

LT2326/LT2926

Assignment 1 - Thai & English OCR

name: Daniilidou Viktoria Paraskevi

How to run the scripts:

1. Splitting the dataset & challenges

In order to get the splitting of the data you have to run on the server this command

```
python splitdata.py 'Thai_English_normal' 'All_Thai_styles' 'Thai_bold_400dpi'  
/scratch/lt2326-2926-h24/ThaiOCR/ThaiOCR-TrainigSet
```

in this path 'Thai_English_normal' is the training, 'All_Thai_styles' is the validation and 'Thai_bold_400dpi' is the test.

Splitting the data was challenging for two reasons. The first one was to be able to separate and call the bold italic from the bold and the second one was the difficulty to ignore and take only one language, when I wanted for example just the Thai bold 200dpi category.

Also, when I was working on this assignment in the beginning I formed my code in a jupyter notebook so in order not to change it in the script the train/validation/test split is hardcoded in the split_train_val_test() function, something that I know it's not efficient and It would be more beneficial to allow language style and dpi to be configurable from the command line. So, the datasets you can choose for training test and validation are those based on the board that was on canvas.

'Thai_normal_alldpi', 'Thai_normal_200dpi', 'Thai_normal_400dpi',
'Thai_bold_alldpi', 'All_Thai_styles', 'Thai_English_normal', 'Thai_bold_400dpi',
'Thai_bold_italic', 'All_Thai_English_styles'.

```
gusdanivi@GU.SE@mltgpu:~$ python splitdata.py 'Thai_English_normal' 'All_Thai_styles' 'Thai_English_normal' /scratch/lt2326-2926-h24/ThaiOCR/ThaiOCR-TrainigSet -h  
usage: splitdata.py [-h] train_dataset val_dataset test_dataset directory  
  
Split data into training, validation, and test datasets.  
  
positional arguments:  
  train_dataset  Dataset for training ( 'Thai_normal_alldpi', 'Thai_normal_200dpi', 'Thai_normal_400dpi', 'Thai_bold_alldpi', 'All_Thai_styles',  
    'Thai_English_normal', 'Thai_bold_400dpi', 'Thai_bold_italic', 'All_Thai_English_styles')  
  val_dataset    Dataset for validation ( 'Thai_normal_alldpi', 'Thai_normal_200dpi', 'Thai_normal_400dpi', 'Thai_bold_alldpi', 'All_Thai_styles',  
    'Thai_English_normal', 'Thai_bold_400dpi', 'Thai_bold_italic', 'All_Thai_English_styles')  
  test_dataset   Dataset for testing ( 'Thai_normal_alldpi', 'Thai_normal_200dpi', 'Thai_normal_400dpi', 'Thai_bold_alldpi', 'All_Thai_styles',  
    'Thai_English_normal', 'Thai_bold_400dpi', 'Thai_bold_italic', 'All_Thai_English_styles')  
  directory      The path of the dataset (e.g., /scratch/lt2326-2926-h24/ThaiOCR/ThaiOCR-TrainigSet).  
  
options:  
  -h, --help      show this help message and exit
```

2. Creating the model & challenges

In order to get the training you have to run on the server this command :

```
python train.py --epochs 10 --batch_size 32
```

and for the test:

```
python test.py --epochs 10 --batch_size 32
```

model.py

I chose in my model to have three layer CNN with increasing filter sizes in order to capture more complex features of the images and I used MaxPooling after each layer so that the spatial dimensions can be reduced. I also used `x.view(x.size(0), -1)` to flatten the tensor for the fully connected layers, so I can get from the multidimensional convolutional layers, a 1-dimensional vector. I added a Dropout layer (0.5) to randomly drop neurons during training, helping the model generalize better and batch normalization to make training more stable.

dataset.py

In this script I have the `extract_unique_labels()` function in order to extract and sort the unique labels from the image paths. Also, inside the `ThaiEngOCRDataset` Class only the .bmp paths are stored and mapped to indices and returns the transformed image and its corresponding label.

train.py

Training the model: In order to ensure that the images have the same format, I used transforms in order to resize the images to 32x32. I created a mapping from labels to indices and `NLLLoss` as a loss function. Before each batch I used `optimizer.zero_grad()` to clear gradients and I compute loss and update the model using backpropagation.

test.py

Evaluates the model using the sklearn metrics accuracy, precision, recall and F1-score on unseen data.

validation.py

Evaluates the Thai OCR model using the sklearn metrics accuracy, precision, recall and F1-score during training.

Checking all the possible character resolutions from the board by running the training we can observe that the loss is getting lower so the model is learning and improving its performance. Also the accuracy, precision, recall and F1-score are high which indicated

that the model performs well. Below you will find screenshots with the training and testing from all the experiments I did with the training and testing data.

A. Training: Thai normal 200dpi, Testing: Thai normal 200dpi

```
gusdanivi@GU.GU.SE@mltgpu:~$ python splitttrial.py 'Thai_normal_200dpi' 'Thai_normal_200dpi' 'Thai_normal_200dpi' /scratch/lt2326-2926-h24/ThaiOCR/ThaiOCR-Tr
ainigSet
Data saved successfully.
gusdanivi@GU.GU.SE@mltgpu:~$ python train.py --train_file training_data.txt --test_file test_data.txt --model_path ./Tha
iEng_model.pth --epochs 10 --batch_size 32
Epoch 1/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:08<00:00, 53.72it/s]
Training Loss after epoch 1: 1.81622114420733
Epoch 2/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 56.70it/s]
Training Loss after epoch 2: 0.6839974471933268
Epoch 3/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:08<00:00, 56.60it/s]
Training Loss after epoch 3: 0.49441765773441093
Epoch 4/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 56.75it/s]
Training Loss after epoch 4: 0.4272011417410363
Epoch 5/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 56.97it/s]
Training Loss after epoch 5: 0.37780302289305934
Epoch 6/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 57.02it/s]
Training Loss after epoch 6: 0.32495328402413964
Epoch 7/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 57.93it/s]
Training Loss after epoch 7: 0.31419921278229873
Epoch 8/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 57.80it/s]
Training Loss after epoch 8: 0.2827999666651614
Epoch 9/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 57.94it/s]
Training Loss after epoch 9: 0.25994327354345625
Epoch 10/10
100%|██████████████████████████████████████████████████████████████████████████████| 453/453 [00:07<00:00, 57.92it/s]
Training Loss after epoch 10: 0.2362374584241945
Model saved to ./ThaiEng_model.pth
model.load_state_dict(torch.load("./ThaiEng_model.pth"))
100%|██████████████████████████████████████████████████████████████████████████████| 57/57 [00:01<00:00, 50.98it/s]
Test Accuracy: 0.967439293598234
Test Precision: 0.9675158761278472
Test Recall: 0.9687634873744693
Test F1-Score: 0.9666131820323207
```

B. Training: Thai normal 200dpi , Testing: Thai normal 400 dpi

```
gusdanivi@GU.GU.SE$ python splitttrial.py 'Thai_normal_400dpi' 'Thai_normal_200dpi' 'Thai_normal_200dpi' /scratch/Lt2326-2926-h24/ThaiOCR/ThaiOCR-Tr  
ainigSet  
Data saved successfully.  
gusdanivi@GU.GU.SE$ python train.py --train_file training_data.txt --test_file test_data.txt --model_path ./ThaiEng_model.pth --epochs 10 --batch_s  
ize 32  
Epoch 1/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 54.70it/s]  
Training Loss after epoch 1: 1.6787679704727807  
Epoch 2/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.54it/s]  
Training Loss after epoch 2: 0.4973008439942816  
Epoch 3/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.46it/s]  
Training Loss after epoch 3: 0.3377833857134363  
Epoch 4/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.72it/s]  
Training Loss after epoch 4: 0.2755001797308547  
Epoch 5/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.47it/s]  
Training Loss after epoch 5: 0.23017288009715686  
Epoch 6/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.58it/s]  
Training Loss after epoch 6: 0.20236530117948667  
Epoch 7/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.66it/s]  
Training Loss after epoch 7: 0.1804686207453318  
Epoch 8/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.13it/s]  
Training Loss after epoch 8: 0.1503489803970204  
Epoch 9/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.73it/s]  
Training Loss after epoch 9: 0.1573422569437404  
Epoch 10/10  
100%|██████████████████████████████████████████████████████████████████████████████| 433/433 [00:07<00:00, 57.59it/s]  
Training Loss after epoch 10: 0.14170062266436445  
Model saved to ./ThaiEng_model.pth
```


F. Training & Testing : Thai and English normal

```
gusdanivi@GU.GU.SE@mltgpu:~$ python splittrial.py 'Thai_English_normal' 'Thai_English_normal' 'Thai_English_normal' /scratch/lt2326-2926-h24/Thai0CR/Thai0CR-TrainigSet
Data saved successfully.
gusdanivi@GU.GU.SE@mltgpu:~$ python train.py --train_file training_data.txt --test_file test_data.txt --model_path ./ThaiEng_model.pth --epochs 10 --batch_size 32
Epoch 1/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:55<00:00, 55.37it/s]
Training Loss after epoch 1: 1.2490413374540146
Epoch 2/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:56<00:00, 53.83it/s]
Training Loss after epoch 2: 0.6765250082432092
Epoch 3/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:54<00:00, 56.15it/s]
Training Loss after epoch 3: 0.5542299120120348
Epoch 4/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:54<00:00, 56.53it/s]
Training Loss after epoch 4: 0.4547522358728064
Epoch 5/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:53<00:00, 57.06it/s]
Training Loss after epoch 5: 0.3967660921061632
Epoch 6/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:53<00:00, 56.99it/s]
Training Loss after epoch 6: 0.3242453776541475
Epoch 7/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:54<00:00, 56.32it/s]
Training Loss after epoch 7: 0.2862937799194424
Epoch 8/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:54<00:00, 56.28it/s]
Training Loss after epoch 8: 0.24822465237489272
Epoch 9/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:53<00:00, 57.10it/s]
Training Loss after epoch 9: 0.2125477762192978
Epoch 10/10
100%|██████████████████████████████████████████████████████████████████████████| 3054/3054 [00:53<00:00, 57.01it/s]
Training Loss after epoch 10: 0.18642150797895213
Model saved to ./ThaiEng_model.pth

model.load_state_dict(torch.load("./ThaiEng_model.pth"))
100%|██████████████████████████████████████████████████████████████████████████| 382/382 [00:06<00:00, 55.22it/s]
Test Accuracy: 0.9623259623259623
Test Precision: 0.9724458429820225
Test Recall: 0.962983676142437
Test F1-Score: 0.963580962640487
```

G. Training & Testing : Thai & English all styles

```
gusdaniv@GU.GU.SE@mlltgu:~$ python splittrial.py --train_file training_data.txt --test_file test_data.txt --model_path ./ThaiEng_model.pth --epochs 10 --batch_size 32  
Epoch 1/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:40<00:00, 55.37it/s]  
Training Loss after epoch 1: 1.0176461825572378  
Epoch 2/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:38<00:00, 56.02it/s]  
Training Loss after epoch 2: 0.47772247045994615  
Epoch 3/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:37<00:00, 56.16it/s]  
Training Loss after epoch 3: 0.2780584805612758  
Epoch 4/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:37<00:00, 56.12it/s]  
Training Loss after epoch 4: 0.1896189590104205  
Epoch 5/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:42<00:00, 55.08it/s]  
Training Loss after epoch 5: 0.15540410318797698  
Epoch 6/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:40<00:00, 55.41it/s]  
Training Loss after epoch 6: 0.13583418015594811  
Epoch 7/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:42<00:00, 55.01it/s]  
Training Loss after epoch 7: 0.1226445441780608  
Epoch 8/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:38<00:00, 56.02it/s]  
Training Loss after epoch 8: 0.11300082417228108  
Epoch 9/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:38<00:00, 56.04it/s]  
Training Loss after epoch 9: 0.10578073446110708  
Epoch 10/10  
100%|██████████████████████████████████████████████████████████████████████████| 12229/12229 [03:38<00:00, 55.90it/s]  
Training Loss after epoch 10: 0.09986329056520338  
Model saved to ./ThaiEng_model.pth
```

```
model.load_state_dict(torch.load("./ThaiEng_model.pth"))  
100%|██████████████████████████████████████████████████████████████████████████| 1529/1529 [00:26<00:00, 56.79it/s]  
Test Accuracy: 0.9699726094599567  
Test Precision: 0.9727352225866035  
Test Recall: 0.9711114082111313  
Test F1-Score: 0.9711634658096616
```